



Automating Oracle database administration, practical examples with Ansible

Mikael Sandström / Ilmar Kerm

@oravirt

@ilmarkerm

Kindred in Numbers

All numbers from Q4 2017

GWR:

£238m

Ebitda:

£75m

Customers:

20m

Daily transactions:

16-18m

GWR from locally regulated markets:

42%

GWR from mobile:

70%

Employees:

1,400

Offices:

15

Nationalities:

52

Brands:

11

Our brands

A portfolio of brands within sportsbook, casino and games





Where we operate



Our Offices

Stockholm | Copenhagen | Tallinn | Antwerp | Sydney
London | Paris | Milan | Gibraltar | Malta | Darwin



Our Licenses

Australia | Belgium | Denmark | Estonia | France | Germany
Gibraltar | Ireland | Italy | Malta | Romania | United Kingdom

WHOAMI: Ilmar Kerm



- Senior database administrator in Kindred Group
- Working in IT since 2000
- Working with Oracle database since 2003
- Working at Kindred since 2007/2015
- From LAMP developer to Oracle DBA
- Oracle ACE

Blog: <https://ilmarkerm.eu>

Email: ilmar.kerm@kindredgroup.com

Twitter: [@ilmarkerm](https://twitter.com/ilmarkerm)



WHOAMI: Mikael Sandström



- Senior database administrator at Kindred Group
- Worked with Oracle for a long time
- Working at Kindred since early 2016
- Automation junkie

Blog: <https://oravirt.wordpress.com>

Email: mikael.sandstrom@kindredgroup.com

Twitter: @oravirt

Ansible for Oracle use cases



- Easy
 - Setting up database environment infrastructure, OS
 - Installation of RDBMS software
 - Automation of routine OS tasks
- Medium
 - Push of the button datacenter deployment
 - Initial creation of databases, pushing out initial configuration
 - Automation of routine DBA tasks: patching, creating tablespaces, ...
 - One time tasks needing high level of orchestration – upgrades, ...
 - Implementing missing features from Oracle
 - Enforcing common security settings
- Advanced
 - All database configuration and changes deployed as version controlled code
 - Easy to review and audit, repeatable

Agenda



- We expect that you are familiar with Ansible
- Things we cover
 - Setting up Ansible environment for Oracle DB use: inventory, modules, cxOracle
 - Handling Ansible inventory for database usage
 - Creating/managing users and roles, syncing with Active Directory
 - Services, tablespaces, resource manager, DBMS_SCHEDULER, AWR settings, INIT.ORA parameters, ASM diskgroups
 - User privileges, including grants to entire target schema
 - Gathering facts about databases
 - Managing passwords in playbooks
 - Practical playbook deployment: execution, scheduling, configuration
- Things we do not cover, because every other Ansible presentation does these
 - What is Ansible (check the first bullet point)
 - Managing OS
 - Installing Oracle RDBMS/Grid software
 - Creating databases/clusters (we will create PDBs though)



Introducing Oracle modules for Ansible

- Modules for administering Oracle database infrastructure
- Hosted in Github
 - <https://github.com/oravirt/ansible-oracle-modules>
- Currently not part of Ansible standard modules
 - Work in progress
- Written in Python
- Require cx_Oracle
 - Needed where the modules are executed
 - cx_Oracle needs an Oracle client
 - If run locally on ‘control-machine’ Instant Client is fine, otherwise just install on the database host
 - pip or yum/rpm

Introducing Oracle modules for Ansible



- Currently manages:
 - Users and roles, including synchronization from AD/LDAP
 - Grants: system and object privileges, including wildcards
 - Tablespaces
 - Parameters
 - Services
 - ASM diskgroups and volumes
 - Pluggable databases
 - DBMS_SCHEDULER jobs, windows and schedules
 - Resource manager consumer groups, including mappings by user profile
 - AWR snapshot settings
- Helper modules:
 - Gathering database facts
 - Executing SQL

Setting up – playbook structure



```
|____ansible.cfg
|____site.yml
|____build-host.yml
|____db-config.yml
|____group_vars
| |____doag      git clone https://github.com/oravirt/ansible-oracle-modules library
| |____doag.yml
| |____passwords.yml
|____library
| |____oracle_asmdg
| |____oracle_awr
| |____oracle_facts
| |____oracle_ldapuser
| |____oracle_parameter
| |____...
|____roles
| |____host
| | |____defaults
| | |____main.yml
| |____tasks
| | |____main.yml
| |____templates
| |____oracle-seclimits.conf.j2
| |____defaults
| | |____main.yml
| |____tasks
| | |____main.yml
| |____oracle-database
| | |____defaults
| | |____main.yml
| | |____tasks
| | |____main.yml
| | |____templates
| | |____dbca.rsp.j2
| | |____netca.rsp.j2
```



Ansible inventory

- Inventory describes for Ansible where the servers are
- Inventory only describes hosts
- Hosts can be divided into groups
- NB! Hosts may not exactly map into databases

Inventory and databases [Potential problem]



- One config file per db

group_vars/**db1**

- Inventory

[**db1**]

host1.example.com

```
ansible-playbook dbconfig.yml -e db=db1
```

- Add another db config (db2), on the same host

group_vars/db1

group_vars/**db2**

- Inventory

[**db1**]

host1.example.com

[**db2**]

host1.example.com

```
ansible-playbook dbconfig.yml -e db=db2
```

- You might get unexpected results (variables might not come from the configuration you expected)



Inventory and databases [Solutions]

- Separate inventory files, one config per inventory

```
$ cat inventory/db1
```

```
[db1]
```

```
host1.example.com
```

```
$ cat inventory/db2
```

```
[db2]
```

```
host1.example.com
```

and target the config with 'ansible-playbook ... -i inventory/**db1**' etc

- Use fake hosts

```
[db1]
```

```
host-db1 ansible_host=host1.example.com
```

```
[db2]
```

```
host-db2 ansible_host=host1.example.com
```

- Use one DNS CNAME per database, same concept as fake hosts
- Use one database per host



Inventory (random thoughts)

- Try to keep the inventories as standard as possible (for as long as possible)
- Maybe keep inventory in its own repository (to avoid duplication, if many different repos for management).
 - Possible to softlink, or just point playbook to inventory
- If environment is large, external service registry could be needed
 - Dynamic inventory script, like for AWS



Setting up scenarios – common variables

```
oracle_base: /u01/app/oracle
oracle_home: "{{ oracle_base }}/product/{{ oracle_db.version }}/db"
oracle_env:
    ORACLE_HOME: "{{ oracle_home }}"
    LD_LIBRARY_PATH: "{{ oracle_home }}/lib"
user: system
oracle_db:
    name: orcl
    passwd: "{{ oracle_passwd }}" #<- Password taken from external file
    version: 12.2.0.1
    cdb: true
    db_dest: +DATA
    fra_dest: +DATA
    pdb:
        - appdb1
        - appdb2
```

Scenario: Setting init.ora parameters 1/2 variables



init_parameters:

- name: db_recovery_file_dest_size
 value: 60G
- name: db_recovery_file_dest
 value: "{{ oracle_db.fra_dest }}"
- name: db_create_file_dest
 value: "{{ oracle_db.db_dest }}"
- name: filesystemio_options
 value: setall
 scope: spfile
- name: optimizer_adaptive_features
 state: absent
- name: audit_trail
 value: XML
 scope: spfile
- name: sec_case_sensitive_logon
 value: TRUE
- name: sec_max_failed_login_attempts
 value: 3
- name: sql92_security
 value: TRUE
 scope: spfile



Scenario: Setting init.ora parameters 2/2 task

```
- name: Set parameters
  local_action:
    module: oracle_parameter
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ oracle_db.name }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
    name: "{{ item.name }}"
    value: "{{ item.value }}"
    state: "{{ item.state | default('present') }}"
    scope: "{{ item.scope | default('both') }}"
  environment: "{{ oracle_env }}"
  with_items: "{{ init_parameters }}"
  tags: param
```



Scenario: Creating PDB-s

```
- name: Create pdb
  local_action:
    module: oracle_pdb
    hostname: "{{ inventory_hostname }}"
    username: sys
    password: "{{ oracle_db.passwd }}"
    service_name: "{{ oracle_db.name }}"
    mode: sysdba
    name: "{{ item }}"
    datafile_dest: "{{ oracle_db.db_dest }}"
    state: present
  environment: "{{ oracle_env }}"
  with_items: "{{ oracle_db.pdb }}"
  tags: pdb
```



Scenario: Creating database users 1/3 variables

`dba_user:`

- `schema: dbauser`
`schema_password_hash: "{ {dbauuser_pwhash} }"`
`default_tablespace: dbauser_data`
`default_temp_tablespace: temp`
`grants:`
 - `dba`

`roles:`

- `name: approle1`
`role_grants:`
 - `create session`
 - `create table`
 - `select any dictionary`

Scenario: Creating database users 2/3 task



```
- name: Create DBA user
  local_action:
    module: oracle_user
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ oracle_db.name }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
    schema: "{{ item.0.schema }}"
    schema_password_hash: "{{ item.0.schema_password_hash }}"
    state: "{{ item.0.state | default('present') }}"
    default_temp_tablespace: "{{ item.0.default_temp_tablespace }}"
    default_tablespace: "{{ item.0.default_tablespace }}"
    grants: "{{ item.1 }}"
    update_password: on_create
  environment: "{{ oracle_env }}"
  with_subelements:
    - "{{ dba_user }}"
    - grants
  tags: users
```



Scenario: Creating database users 3/3 task

```
- name: Create application role
  local_action:
    module: oracle_role
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ oracle_db.name }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
    role: "{{ item.name }}"
    state: present
  environment: "{{ oracle_env }}"
  with_items: "{{ roles }}"
  tags: role

- name: Add grants to role
  local_action:
    module: oracle_grants
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ oracle_db.name }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
    role: "{{ item.name }}"
    grants: "{{ item.role_grants }}"
    state: present
  environment: "{{ oracle_env }}"
  with_items: "{{ roles }}"
  tags: role,grant
```

Scenario: Syncing users and roles from AD



```
- name: oracle_ldapuser
  local_action:
    module: oracle_ldapuser
    ##### Target Oracle database connection information is cut from this slide
    ##### LDAP parameters
    ldap_connect: "ldap://mycompany.domain:389"
    ldap_binddn: Administrator@mycompany.domain
    ldap_bindpassword: Ansible2017
    ldap_user_basedn: OU=Users,OU=Demo,DC=ansible,DC=test
    # The following filter means that objectClass is person, member of one specific group and account is not disabled
    ldap_user_filter: (&(objectClass=person)(memberOf=CN=prod_db,OU=Groups,DC=ansible,DC=test)(!(userAccountControl:1.2.840.113556.1.4.803:=2)))
    ldap_username_attribute: sAMAccountName
    ##### New user parameters
    user_profile: LDAP_USER
    user_grants:
      - create session
      - create table
    group_role_map:
      - {dn: "CN=HR,OU=Groups,OU=Demo,DC=ansible,DC=test", group: "department_hr"}
      - {dn: "CN=IT,OU=Groups,OU=Demo,DC=ansible,DC=test", group: "department_it"}
    deleted_user_mode: lock
  environment: "{{ oracle_env }}"
```

All synced users must have a common dedicated profile!

Users missing from LDAP query will be locked or deleted.

Scenario: Maintaining wildcard privileges



```
- name: Manage what objects DEPARTMENT_HR can access
  local_action:
    module: oracle_privs:
      ### Target Oracle database connection information is cut from this slide
      roles:
        - DEPARTMENT_HR
      privs:
        - SELECT
        - FLASHBACK
      objs:
        - HR.%
        - HRREST.COUNTRIES
      objtypes:
        - TABLE
        - VIEW
      state: present
  environment: "{{ oracle_env }}"
```

```
SELECT owner, object_name
FROM all_objects
WHERE (owner||'.'||object_name LIKE
'HR.%' OR owner||'.'||object_name =
'HRREST.COUNTRIES')
AND object_type IN ('TABLE','VIEW')
```

Existing extra privileges not matching the privs + objects list will be revoked!

Scenario: Managing resource manager



```
- name: Manage consumer group for LDAP synced users
  local_action:
    module: oracle_rsrc_consgroup
    ### Target Oracle database connection information is cut from this slide
    name: ad_user_consgroup
    state: present
    # NB! Oracle resource manager actually does not support
    # user profile based mapping, so we expand it to user list
    map_oracle_user_profile: LDAP_USER
    grant_user_profile: LDAP_USER
  environment: "{{ oracle_env }}"
```



Managing passwords

- You should not store plain text passwords in playbooks!
- Ansible-vault
 - Ansible native solution to encrypt YML files that can be included in playbooks
 - Ansible-vault files can be committed to git
 - Password delivered through user input, file or script
- Oracle wallet
 - Ansible-oracle-modules use cx_Oracle python module that requires OCI client, therefore it also supports Oracle wallet authentication
 - Requires setting up tnsnames.ora and sqlnet.ora
 - Each TNS name needs password entry in wallet
 - Limitation for large deployments?
- Calling remote password service from playbook
 - REST API calls to e.g Passwordstate to fetch individual DB passwords
 - API key can be secured with ansible-vault

Scenario: Securing passwords with Ansible Vault



```
$ ansible-vault create passwords.yml
```

```
sys_password: Oracle123
```

```
dbauser_password: Oracle456
```

```
$ cat passwords.yml
```

```
$ANSIBLE_VAULT;1.1;AES256
```

```
3730663139323136333439666643466613062373336623836383333653963376463366234353130  
3962383634356134663061313162386136386234616138380a616630626535613131353431356536  
6662383131663730643236376333323132313936353566613766633335346363376363333383963  
3535386137633165660a37643561373666663396233626531346530343462333865306137386434  
35303466646564313131386565386364333762646130316439343764356366643034623563633061  
32373832376136303834356664373534643631303439656434613761633732653433646536346231  
343062396162376366326561333537343736
```

Scenario: Create job in all PDB 1/3 variables



```
adb_json_query: adb[?OPEN_MODE=='READ WRITE'].NAME
```

```
exclude_pdb:
```

- PDB\$SEED

```
common_job:
```

```
  name: DBAUSER.SEND_SLOW_SQL_REPORT
```

```
  job_type: plsql_block
```

```
  job_action: DBA_REPORTS_PKG SEND_SLOW_SQL_REPORT
```

```
  repeat_interval: FREQ=DAILY;BYHOUR=8
```

Scenario: Create job in all PDB 2/3 tasks



```
# Connect to CDB and gather facts
- name: Gather target CDB facts
  local_action:
    module: oracle_facts
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ oracle_db.name }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
  register: dbfacts
  environment: "{{ oracle_env }}
```

Scenario: Create job in all PDB 3/3 tasks



```
# Connect to all PDB-s and create the job
- name: Create a small test job in all PDBs
  local_action:
    module: oracle_job
    hostname: "{{ inventory_hostname }}"
    service_name: "{{ item }}"
    user: "{{ oracle_db.user }}"
    password: "{{ oracle_db.passwd }}"
    # Job info
    state: present
    name: "{{ common_job.name }}"
    job_type: "{{ common_job.job_type }}"
    job_action: "{{ common_job.job_action }}"
    repeat_interval: "{{ common_job.repeat_interval }}"
  environment: "{{ oracle_env }}"
  when: dbfacts.ansible_facts.cdb
  with_items: "{{ dbfacts.ansible_facts | json_query(pdb_json_query) | difference(exclude_pdb) }}"
```

Task only runs when target database is CDB. Task is repeated for each PDB.



Scheduling/running playbooks

- We use Jenkins
- Basically 3 ways of running jobs
 - Manually
 - Trigger based
 - File is committed to Git
 - Git HEAD branch is merged to stable PROD branch
 - Scheduled



Running playbooks (a few other options)

- Ansible Tower (commercial offering)
- Ansible AWX (open source Tower)
- Rundeck
- Semaphore
- *Tensor*
- *Polemarch*



Oracle infrastructure as code

- No changes should be done manually on database infrastructure
- Changes done in Ansible code only
- Ansible code version controlled in git
- Commit/merge to git automatically executes Ansible playbooks that deploy the change



DEMO: Oracle infrastructure as code

- Demo environment:
 - Oracle 12cR2, with one CDB
 - Git
 - Bitbucket
 - Ansible executed from Jenkins
- Action flow:
 - Change in playbook committed to git (Bitbucket)
 - Add new ASM diskgroup
 - Change init.ora parameters in CDB
 - Create and initialise new application PDB
 - Add role with system grants
 - Add tablespaces
 - Add application schema
 - Add service
 - Bitbucket commit hook fires a job in Jenkins



Ansible is great, but...

- ... it is not for everything
- There are better solutions for:
 - Managing schema objects
 - Deploying database code
- Better look at specific tools for these purposes:
 - Liquibase
 - FlyWay
- But... Ansible can execute these tools as part of the release process orchestration



Resources

- Ansible homepage: <https://www.ansible.com/>
- Oracle database modules for Ansible:
<https://github.com/oravirt/ansible-oracle-modules>
- Deploying RAC with Ansible:
<https://github.com/oravirt/ansible-oracle>





Download scripts

- Here you can download the scripts used in this presentation
 - <https://github.com/oravirt/oug-ansible-oracle>





Thank You